

UniFLEX™ Sort/Merge



technical systems
consultants, inc.

UniFLEXTM **Sort/Merge**

COPYRIGHT © 1981 by
Technical Systems Consultants, Inc.
111 Providence Road
Chapel Hill, North Carolina 27514
All Rights Reserved

TM UniFLEX is a trademark of Technical Systems Consultants, Inc.

MANUAL REVISION HISTORY

| Revision | Date | Change |
|----------|------|------------------|
| A | 2/81 | Original Release |

COPYRIGHT INFORMATION

This entire manual is provided for the personal use and enjoyment of the purchaser. Its contents are copyrighted by Technical Systems Consultants, Inc., and reproduction, in whole or in part, by any means is prohibited. Use of this program, or any part thereof, for any purpose other than single end use by the purchaser is prohibited.

DISCLAIMER

The supplied software is intended for use only as described in this manual. Use of undocumented features or parameters may cause unpredictable results for which Technical Systems Consultants, Inc. cannot assume responsibility. Although every effort has been made to make the supplied software and its documentation as accurate and functional as possible, Technical Systems Consultants, Inc. will not assume responsibility for any damages incurred or generated by such material. Technical Systems Consultants, Inc. reserves the right to make changes in such material at any time without notice.

TABLE of CONTENTS

| | Page |
|--|------|
| 1.0 Introduction | 1 |
| 2.0 Sorting Concepts | 3 |
| 3.0 Sort Tutorial | 7 |
| 3.1 Sample Sort | 8 |
| 3.2 Sorting with input keys | 11 |
| 3.3 Sorting with output keys | 13 |
| 3.4 Sorting with fields | 16 |
| 3.5 Sorting with right or left justification | 20 |
| 4.0 General User Specifications and Sort/Merge Details | 21 |
| 4.1 Specifying input records | 21 |
| 4.2 Specifying input keys | 22 |
| 4.3 Specifying output source | 23 |
| 4.4 Specifying output keys | 24 |
| 4.5 Specifying output redirection | 25 |
| 4.6 Specifying input redirection | 26 |
| 4.7 Location of the sort/merge module | 26 |
| 4.8 Temporary files | 27 |
| 5.0 The SORT Command | 29 |
| 5.1 General use | 29 |
| 5.2 Additional parameter options | 32 |
| 5.3 Select/Exclude option | 34 |
| 6.0 The PSORT Command | 37 |
| 7.0 The CSORT Command | 39 |
| 8.0 The MERGE Command | 43 |
| 9.0 The PMERGE Command | 45 |
| 10.0 Run-Time Messages | 47 |
| 11.0 Error Messages | 49 |
| 11.1 From SORT and MERGE | 49 |
| 11.2 From CSORT | 50 |
| 11.3 From PSORT and PMERGE | 52 |
| 12.0 Alternate Collating Sequence File | 55 |
| 13.0 Tag File Output | 57 |
| 14.0 Parameter File Description | 59 |
| 14.1 Parameter file layout | 59 |
| 14.2 Calling Sort/Merge from other programs | 63 |
| 15.0 Appendix A - ASCII Table | 65 |

UniFLEX Sort/Merge User's Manual

1.0 INTRODUCTION to the UniFLEX™ SORT/MERGE PACKAGE

The UniFLEX Sort/Merge Package is a very powerful and quite complex program. It was, however, designed with the operator in mind and is relatively simple to use. This manual was written with the non-computerist in mind and will lead you gradually into use of the sort/merge package. The user is advised to experiment with the use of the sort/merge on the sample data file supplied on the disk or with a non-critical file of his own. The best way to learn the operation of the software is from hands-on experience.

As the name implies, the sort/merge package has two major functions, sorting and merging. The most used function will be that of sorting and will be the first discussed in this manual. Much of the description of the sort, however, will also apply to the merge.

In order to use the UniFLEX Sort/Merge Package, you must have a functioning UniFLEX system and should be somewhat familiar with its operation. The system manager should copy the five command files named "sort", "csort", "psort", "merge", and "pmerge" from the UniFLEX Sort/Merge distribution disk into the "/bin" or "/usr/bin" directory of the UniFLEX system. The files "names" and "names2" are sample text files which are referred to throughout this manual. They may be copied into an appropriate directory for access by users if desired.

The UniFLEX Sort/Merge Package permits a user to sort the contents of almost any file into almost any desired order. Output may be routed to terminal, disk file, or most any other device.

™ UniFLEX is a trademark of Technical Systems Consultants, Inc.

2.0 SORTING CONCEPTS

The UniFLEX Sort/Merge Package can be used to sort almost anything which is contained in a disk file. Generally it is used to sort some type of textual data such as a file of names and addresses. The "items" or pieces of information which are sorted into order are called "records". If, for example, you had a file containing a list of 10 peoples last names, each name would be called a record. If the list contained the last name followed by the first name, then one record would include the last name and the first name. The sort routine may be used to rearrange these records into some defined order. For example you could sort the list of last names into alphabetical order or perhaps into reverse alphabetical order (Z's at the top of the list). If sorting the file described above which had first and last names, it would be possible to sort according to the last name only so that the finished output would have last names in alphabetical order with the first names following along. Alternatively, we could sort according to first name, so that the first names would be in order with the last names following along. In short, it is possible to sort according to any portion of the input record. This portion which is used as a reference for sorting is called an "input key". Thus we could specify to the sort package that we wanted the last name as the key or that we wanted the first name. We could go a step further and specify more than one key. That is, we might sort first on the last name and then on the first name. This would mean that anytime there were several records with identical last names, they would end up being sorted together with the corresponding first names also put into order. Or we could specify the first name as the first input key and the last name as the second input key. This would sort the records into order by first name and where there were multiple records with identical first names, the corresponding last names would also then be put into order. If the file we wished to sort had more than just a first and last name (such as address, phone number, etc.), we could sort on several different keys. It is possible to specify as many as 20 input keys to the sort routine. Each key may be individually specified as ascending or descending (ie. alphabetical or reverse alphabetical), right or left justified (trailing or leading blanks omitted), and may be specified as any portion of the input record by specifying column numbers. For example, we might want to sort on columns 1 through 10 (inclusive) which we know contains a last name in each record. Or perhaps on columns 21 through 30 (inclusive) which is known to contain a phone number in each record.

In the simplest case, the output file is a rearranged version of the records of the input file. Note that it is not necessary to actually produce a disk file of the output. If desired, the output can be routed to the terminal or a printer. When a disk output file is produced, it is entirely independent of the input file. In other words, the input file is not altered in any way. The sort simply reads through the input file and produces a new file with the output. If desired, the output records (one "record" of information is sent to the output for each record which is input) do not have to be a carbon copy of the input. You may specify that only certain portions of the input record be sent

to the output record. These "portions" are called output keys much like the input keys described before. Besides allowing you to output portions of the input file, they allow tabbing to some specific column number on the line being output and allow you to insert some comment or characters into each output record.

Another option which is supported is that of an "alternate collating sequence file". Sorting is normally done according to the ASCII sequence of characters. ASCII stands for American Standard Code for Information Interchange which is simply a widely used code or sequence of characters. See Appendix A for a chart of the ASCII sequence. If it is desirable in some instance to sort according to some other sequence of characters, that sequence may be supplied to the sort package in the form of a disk file containing the characters in the desired order. This alternate collating sequence file will then be used instead of ASCII for sorting purposes.

Other features of the sort/merge package include the ability to treat lower case characters equivalent to upper case characters for sorting purposes. This means that "Jones", "JONES", and "jones" would all be considered equal. It is possible to have the sort program automatically delete or ignore any records which have an entirely blank (all spaces) sort key, where "sort key" refers to all the input keys put together. There is also a "select/exclude" option which allows input records to be selected or excluded from the sort depending on their contents.

The sort/merge package is so called for two reasons. One is that it is capable of merging two or more input files which have been previously sorted or are known to be in order. This is called a "merge-only" operation since no actual sorting is done and still requires most of the specifications which the sort requires such as input and output keys. The second reason for being called a sort/merge package is in the way sorting is sometimes done. The actual sorting is always done entirely within the computer's memory. The input records themselves are not sorted, but rather only the input keys. Each key has a pointer to where the input record it came from resides so that when the keys have been sorted, the input records can now be read back in the same order as the keys. This usually saves memory in that the sort key is almost always shorter than the input record. The program reads an input record into a temporary buffer and produces the sort key from the input keys specified. This sort key is then written into the available memory space. This process is repeated until enough sort keys have been written to fill the available memory space or until all records of the input file have been read. If the sort keys for the entire input file all fit in memory, the sort is carried out and the output file is written out directly. If there is not enough memory for the entire input file, however, the program sorts what did fit and then writes that sorted portion of the file out to a temporary work file on disk or "run" as we will often refer to it. Then the input and sort process is repeated on the next section of the input file which will fit in memory and another run is written out. This continues until the entire input file has been read. We now have several temporary work files each of which contains a sorted portion of the original file. The sort/merge package will now merge those work files or runs into one sorted output

file and delete all the temporary files on completion. If there are a large number of runs, it may not be possible to merge them all at once. If this is the case, as many as possible are merged into another single run or work file which may then be merged with the remaining runs in a second pass. In extreme cases (caused by extremely large input files) several of these merge passes may be required.

It should be obvious that the sort/merge package requires a good deal of information or parameters specified in order to know just how to carry out the sort. These parameters are passed to the main sort/merge program by placing them in a disk file called the parameter file which can be read by the sort/merge program. This parameter file can be prepared and/or passed to the main sort/merge program in any one of five ways. The UniFLEX Sort/Merge Package actually consists of five command files. Each of these five commands employs a different method of producing a parameter file and passing it to the main sort/merge program. One of those five files, called "psort", is itself the main sort/merge module and is responsible for the actual sorting and merging processes. The others are concerned only with supplying all the necessary parameters to the "psort" module. When the "psort" module is called, it expects to find all the necessary parameters in a disk file called the "parameter file". "psort" is itself a command and if executed, it reads the specified parameter file and performs a sort/merge according to the parameters in that file. The other four commands prepare a parameter file and then automatically load and execute the "psort" module to carry out the actual sort/merge. Each of the commands has a different way of setting up the parameters as explained below.

The first is called "sort" and is used to set up parameters for a sort operation. It does so by prompting the user for all the necessary information. When all the parameters are set, there is an option for saving them in a disk file so that the very same parameters may be used again (as explained later) without having to re-enter them. If not saved, they will be written to a temporary file and deleted upon completion of the sort.

The second is called "psort" which stands for Parameter file Sort and is also used for a sort operation. This command allows the user to specify a parameter file (which was produced by the "sort" command) by name so that previously setup parameters may be easily recalled. As described above, this is also the actual sort/merge module which will be called by the other four commands when they have prepared the necessary parameter file.

The third is called "csort" which stands for Command line Sort and is also used to setup parameters for a sort. Instead of being prompted for the parameters, the user must supply them in a somewhat condensed form on a single command line. This is quicker and often more convenient for the user who fully understands the operation and use of the sort/merge.

UniFLEX Sort/Merge User's Manual

The fourth is called "merge" and is used to setup parameters for a merge-only operation. The program prompts the user for all the information much as in the "sort" command and then proceeds with the merge. There is also a provision for saving the parameters in a disk file for later use.

The fifth and final command is "pmerge" standing for Parameter file Merge and is also used to prepare for a merge-only operation. Much like "psort", it allows the user to specify a parameter file which has been prepared by either "sort" or "merge".

This modularity gives the user a great variety of capabilities and allows convenient operation of the sort/merge package under almost all circumstances.

3.0 SORT TUTORIAL

The "sort" command is probably the easiest method of using the sort/merge package, especially for the novice. It prompts the user for all the necessary information thus obviating the need to memorize what parameters must be supplied. Most of the information asked for will default to some preset default value if the operator simply hits a carriage return in response to a particular prompt. The prompts usually include what choices the operator has for a response. The choice which will be used as a default is generally designated by a following asterisk ('*'). For example, one prompt you will receive will be

Save parameter file on disk (y or n*)?

The two choices you have as a response are 'y' and 'n' for yes and no. You could, however, simply hit a carriage return and the "sort" program will default to not saving the parameter file since the 'n' response is flagged with an asterisk.

Prompts from "sort" require responses that are terminated with a carriage return. The response may be some typed information followed by a carriage return or if there is no response (you are selecting the default) the user may type a carriage return only. In any event, nothing is done by "sort" until the entire line has been entered and a carriage return typed. When entering this type of response, there are four control characters which may be used as a convenience. Control characters are entered by typing the 'CTRL' key at the same time as some letter key is hit. The possible control keys are CTRL H, CTRL X, CTRL Z, and CTRL D. The CTRL H causes a single backspace in the line being typed. If your keyboard has a backspace key, it will accomplish the same function. The CTRL X is a cancel function which will cancel the line presently being typed and permit entry of a replacement line. The CTRL Z is a restart function meaning that anytime a CTRL Z is hit, the "sort" program will start over with the first prompt. The CTRL D causes "sort" to immediately terminate and returns control back to the operating system. As always under UniFLEX, a CTRL C may also be typed at any time to terminate the running program and return control back to the operating system.

To help clarify the following description of use of the "sort" command, references will be made to a sample data file named, "names" as found on the disk on which the sort/merge package is distributed. It is listed here as a reference.

UniFLEX Sort/Merge User's Manual

| | | |
|---------|----------|---------|
| WILSON | JIM | 4578339 |
| SMITH | JERRY | 4226002 |
| ABBOTT | JIM | 4572091 |
| CASWELL | AMY | 9258411 |
| PERRY | ARNOLD | 4226008 |
| SOUTH | MARGARET | 4221267 |
| SMITH | HAROLD | 4227264 |
| LYON | WILLIAM | 8538227 |
| TYLER | HENRY | 4575573 |
| FREEMAN | LINDA | 4229105 |

There are several things which should be noted about this file. First of all, this is a list of names and phone numbers. Each line of the file should be considered one input record. Note that all records in this particular file are of the same length, namely 27 characters. Each last name begins in column 1 and may run through column 10. Each first name begins in column 11 and may run through column 20. Each phone number is in columns 21 through 27. Most important, however, is the fact that this file is in no particular order!

3.1 SAMPLE SORT

Suppose we wished to sort the file "names" into order according to last names. Let us go through an example of such a sort. At this point, it is not crucial to completely understand all the prompts and responses. Simply go through the steps as instructed and watch the results. More details will be given later. The first step is to be certain that the UniFLEX Sort/Merge Package has been placed into the file system you are operating with and that the sample file "names" is also available in your current directory. Now type the command

```
++ sort names
```

This tells the computer to load the "sort" command and use it to sort the input file named "names". Note that the plus signs are the UniFLEX prompt and are not typed by the user. If the "names" file is not in your current directory, you will have to specify the directory path name along with it. The computer should respond with

```
=== UniFLEX Sort Parameter Editor ===
```

```
Fixed or variable length records (f or v*)?
```

We are sorting variable length records (even though they are all the same length in this case) so type a 'v' or a carriage return. More information on input record type is given later.

The computer will then ask

```
EOR character or field count (default is EOR=$OD)?
```

where EOR stands for End Of Record. Our file has an EOR character which is a carriage return (a hex 0D) so we can either hit a carriage return allowing the default to be set or type a "\$0D" to set the EOR to a carriage return.

Now we are asked

Field separator character?

We do not have "fields" (these will be described later), so simply type a carriage return which defaults to a null field separator character.

Next the computer asks

Output from key, input, or tag (k, i*, or t)?

This prompt gives us the option of having the data in the output records come from the actual input records, from the sort key itself, or from the "tags" as will be described later. In most cases you will probably want the output to come from the input records. That is the case here since we simply want a rearranged version of the input records. Your response should therefore be an 'i' or use the default by typing a carriage return.

Now we get to the real meat of the sort parameters, the input keys. We are prompted with

Enter input keys. Defaults to "a(1)1-10".
?

As mentioned before, the input keys are specified by the starting and ending column numbers of the portion of the input record that is to be used as a sorting reference. Therefore, a part of EVERY input key specification must be of the form "sss-eee", where "sss" represents the starting column and "eee" represents the ending column inclusive. Now in our sample we want to sort on the last names. We can see that they all start in column one and end with column 10. Therefore, our input key specification would simply be "1-10". Type those 4 characters after the question mark. Note that we could use the default input key by typing a carriage return since the "sort" routine will default to "a(1)1-10". The "a(1)" portion of this default specification is some added information about the input key which will be covered later.

When you have entered the "1-10" and hit a carriage return, another question mark ('?') prompt will be issued. This is to allow another input key to be entered if desired. It is not required in our sample since we are sorting only on the last name, so type a carriage return to exit the input key prompting. The computer will respond with

Enter output keys. Defaults to entire record.
?

At this point, you could specify that you only wanted certain portions of the input record sent to the output record. We, however, want the

UniFLEX Sort/Merge User's Manual

entire input record sent as is to the output record. There are several ways to specify this, but the easiest is to use the default by typing a carriage return.

Now we see the prompt

Further options required (y or n*)?

If a 'y' is typed, "sort" will begin prompting for several other possible input parameters. We will not be needing any of these somewhat specialized parameters, so type a 'n' or carriage return.

At this point, "sort" will type

=== Parameters are now set ===

Save parameter file on disk (y or n*)?

This tells us that we have completed the setting up of the sort parameters. If desired, all these parameters may be saved in a disk file so that we could later repeat the sort without needing to retype them. For now, just type an 'n' or a carriage return since the parameters we are using are quite simple to enter.

Finally we see the prompt

Exit or proceed with sort (e or s)?

This allows us to exit back to the disk operating system or to proceed with the sort as specified. Since we did not save our parameters in a file, we want to continue with the sort. Type an 's' to do so. Note that there is no default value in this case. It is necessary to actually type an 'e' or an 's'.

At this point, the program "psort" (the actual sorting module) is loaded and the sort begins. You should see the following information on your terminal.

| | | |
|---------|----------|---------|
| ABBOTT | JIM | 4572091 |
| CASWELL | AMY | 9258411 |
| FREEMAN | LINDA | 4229105 |
| LYON | WILLIAM | 8538227 |
| PERRY | ARNOLD | 4226008 |
| SMITH | JERRY | 4226002 |
| SMITH | HAROLD | 4227264 |
| SOUTH | MARGARET | 4221267 |
| TYLER | HENRY | 4575573 |
| WILSON | JIM | 4578339 |

As you see, the records have been put into ascending alphabetical order by last name. Now if we had wanted, we could have specified that the sorted output be sent to a disk file. This would be done by redirecting the i/o on the command line as described in the UniFLEX operating system documentation.

3.2 SORTING WITH INPUT KEYS

If you will examine just what prompts you were required to fully answer in the preceding example and which ones you were able to use the default value on, you will notice that all responses could be defaulted. Often most responses can be defaulted except for the input keys. Generally your specific sort needs will require you to enter some key other than "1-10". At this point you might try a couple of more sorts on the file "names". Instead of "1-10" for an input key, try sorting by first name only. This could be done by responding in the same manner as the example above to all the prompts except the one for entering input keys. To this you should respond with "11-20" since the first names all start in column 11 and can extend up to column 20. Then try sorting by phone number in the same manner. This would require an input key of "21-27".

If you will examine the output of our first example above, you will see that there are two Mr. Smiths. Their last names are sorted into order as requested, but their first names (which were simply "carried along" during the sort) are not. If we wanted the file sorted first by last name and then by first, there are a couple of ways we could go about it. The simplest and most logical method in this particular case would be to specify a single input key that would include both names since they are placed in the record with last name first. The key would simply be "1-20". However, for instructional purposes, we will perform the sort with multiple input keys. In other words, we will first specify a key for the last name and then another for the first name. The prompt and response would look like this

```
Enter input keys. Defaults to "a(1)1-10".
? 1-10,11-20
?
```

Note that both keys were specified on one line, separated by a comma. In all the examples given in this manual, commas will be used as delimiters, but if desired a space may always be used instead. It would also be possible to put the keys on more than one line as shown

```
Enter input keys. Defaults to "a(1)1-10".
? 1-10
? 11-20
?
```

When you have entered all the keys you wish, you can exit the key input mode (question mark prompts) by hitting a carriage return in response to the question mark.

UniFLEX Sort/Merge User's Manual

If a sort is performed with these input keys and all other responses as before, you should see the following.

| | | |
|---------|----------|---------|
| ABBOTT | JIM | 4572091 |
| CASWELL | AMY | 9258411 |
| FREEMAN | LINDA | 4229105 |
| LYON | WILLIAM | 8538227 |
| PERRY | ARNOLD | 4226008 |
| SMITH | HAROLD | 4227264 |
| SMITH | JERRY | 4226002 |
| SOUTH | MARGARET | 4221267 |
| TYLER | HENRY | 4575573 |
| WILSON | JIM | 4578339 |

Notice that the Smiths are now in order by last AND first name. Experiment at this time with other combinations of multiple input key sorts. For example try sorting on phone number, then last name, then first name. Or perhaps on first name, then last name, then phone number.

Multiple input keys can be entered in any order. That is to say, the first key might come from 21-27 while the second comes from 1-10. It is also possible to use the same portion of the input record more than once. This does not often make sense, but you could, for example, specify a key like "1-8,30-40,1-10".

To this point, all our sorts have been in ascending order. That is 'A' to 'Z' or '0' to '9'. Sometimes you may wish to sort in a descending order. This is very easily accomplished through the manner in which the input keys are specified. Any input key may be prefaced with the letter 'a' for ascending or 'd' for descending. If neither is specified, the key is sorted in ascending order or in other words, 'a' is the default sorting order. That is what is meant by the 'a' in the input key prompt default. As an example, let's sort "names" on the last and first name, but in DESCENDING order. All responses to the "sort" command's prompts are as before (may be defaulted) except for the input key specification. This should now be as shown.

```
Enter input keys. Defaults to "a(1)1-10".
? d1-20
?
```

Performing the sort with these parameters should give you

| | | |
|---------|----------|---------|
| WILSON | JIM | 4578339 |
| TYLER | HENRY | 4575573 |
| SOUTH | MARGARET | 4221267 |
| SMITH | JERRY | 4226002 |
| SMITH | HAROLD | 4227264 |
| PERRY | ARNOLD | 4226008 |
| LYON | WILLIAM | 8538227 |
| FREEMAN | LINDA | 4229105 |
| CASWELL | AMY | 9258411 |
| ABBOTT | JIM | 4572091 |

It is also possible to have multiple input keys with some keys specified in ascending order and some in descending order. If for some reason we wanted the last names sorted in descending order but the first names in ascending order, we could specify input keys of "d1-10,all-20" or taking advantage of the default sort order, "d1-10,11-20".

3.3 SORTING WITH OUTPUT KEYS

We now turn our attention to output key specification. In all the above examples, we output the entire input record by default on the prompt to enter output keys. It is possible to do a certain amount of output formatting by use of output keys. These output keys are similar to input keys in that they allow you to send selected portions of the input record to the sorted output records. As with the input key, the simplest form of output key would be "sss-eee" where 'sss' represents the starting column and 'eee' represents the ending column.

Let's assume we wish to sort the file "names" according to last name, but want the output to be only the last names. You would begin exactly as before with an input key of "1-10", but when prompted with

Enter output keys. Defaults to entire record.
?

do not use the default. Instead, type in the output key necessary to send only the last name to the output record. In this case it would be columns 1 through 10, so enter "1-10". Performing the sort with these parameters would yield the following output.

UniFLEX Sort/Merge User's Manual

ABBOTT
CASWELL
FREEMAN
LYON
PERRY
SMITH
SMITH
SOUTH
TYLER
WILSON

Now try sorting the "names" file on any input keys you desire, but only output the first name or phone number. You can also specify more than one output key. You might try outputting the first name, last name, and then phone number. This would require a list of output keys like "11-20,1-10,21-27".

There is a special symbol which may be used in place of the ending column number in an output key. That is an asterisk ('*') and it represents the end of the input record. In other words, an output key of "11-*" would start outputting from column 11 and continue to the end of the record. For example, "11-20,1-10,21-*" would output first name, last name, and then phone number exactly as before. A "1-*" would mean to output the entire input record. This would have the same effect as defaulting to the entire input record.

It does not make sense to use an 'a' or 'd' (ascending or descending) as a preface to an output key, but there is another type of preface that can be used. That is an 'r' or 'l' which stand for Right or Left justify respectively. Any output key which is prefaced by one of these characters will be either right or left justified. This simply means that in the case of a right-hand justify any blanks or spaces at the end of the key will be moved to the front or in the case of a left-hand justify any leading blanks will be moved to the end of the key. An example will help clarify. Perform a sort on "names" with an input key of "11-20" and respond to the output key prompt as shown

```
Enter output keys. Defaults to entire record.  
? r11-20  
?
```

When the sort is initiated, you should see the following output.

```

      AMY
    ARNOLD
    HAROLD
      HENRY
      JERRY
        JIM
        JIM
      LINDA
    MARGARET
    WILLIAM

```

Notice that the first names are all right justified. This is a convenient feature for producing readable output lists from the sort/merge package. Right or left justify can also be used with input keys as will be described later.

There are three special types of output keys that may be used. One is a simple tab function for tabbing to a specified column and the other two allow the insertion of some string of characters or words that are not in the input record.

The tab function is called by typing an at-sign ('@') followed by the column number (largest allowed = 255) to which you wish to tab. For example, "@25" is a valid output key which will cause a tab to column 25 before any further outputting. The tab is performed by simply inserting the required number of spaces in the output record to reach the specified column.

Another of the three special output key types allows you to specify some particular character to be inserted into the output key. The character is specified by a hexadecimal ASCII value (see Appendix A) preceded by a dollar-sign ('\$'). For example to send a period to the output record we could use a key of "\$2E". This also allows the insertion of control characters into the output record. A control character is a non-printing character whose ASCII value is less than hex 20.

The third special type of output key permits whole strings of characters to be inserted. The string of characters is preceded and followed by an apostrophe or single-quote. For example, a key of 'SAMPLE' (apostrophes included) would insert the word SAMPLE into each output record. Note that an apostrophe cannot be included in the string of characters as it would appear as the string terminator. If an apostrophe is desired in the output key, it must be specified as a hex value character as shown above.

Let's put some of these special output keys to work in another example with our "names" file. Answer all prompts as before until you reach the

UniFLEX Sort/Merge User's Manual

input key prompt. Here use an input key of "1-20" for an alphabetical listing of last and first names. Respond to the output key prompt as shown.

```
Enter output keys. Defaults to entire record.
? r11-20,$20,1-10,'    phone no. ',21-23,'-',24-*
?
```

This looks rather complicated, but if you examine each key one at a time it is really quite simple. We are printing the first name right justified, followed by a space (a hex 20), followed by the last name. Next we print a string which spaces over a few columns, prints "phone no.", and then spaces one more column. Next we output the first three characters of the phone number, a dash, and the rest of the phone number. The result is as follows.

| | |
|----------------|--------------------|
| JIM ABBOTT | phone no. 457-2091 |
| AMY CASWELL | phone no. 925-8411 |
| LINDA FREEMAN | phone no. 422-9105 |
| WILLIAM LYON | phone no. 853-8227 |
| ARNOLD PERRY | phone no. 422-6008 |
| HAROLD SMITH | phone no. 422-7264 |
| JERRY SMITH | phone no. 422-6002 |
| MARGARET SOUTH | phone no. 422-1267 |
| HENRY TYLER | phone no. 457-5573 |
| JIM WILSON | phone no. 457-8339 |

This should give you some idea of the endless possibilities for output records.

3.4 SORTING WITH FIELDS

Now let us examine another feature of the sort/merge package, "fields". Often it is desirable to divide a single input record into a number of distinct portions. We call these portions "fields". In our example with the file "names" we had a last name, first name, and phone number but they were actually not separate fields since each item was separated from the others only by its columnar position in the record. To have fields, there must be some sort of field separator character to split the input record into distinct fields. There is another file on the disk called "names2", which does have fields. It looks like this.

WILSON,JIM,4578339
 SMITH,JERRY,4226002
 ABBOTT,JIM,4572091
 CASWELL,AMY,9258411
 PERRY,ARNOLD,4226008
 SOUTH,MARGARET,4221267
 SMITH,HAROLD,4227264
 LYON,WILLIAM,8538227
 TYLER,HENRY,4575573
 FREEMAN,LINDA,4229105

Notice that the list is exactly as before except that instead of finding the names and number in fixed columns they are packed together with only a comma separating them. These can be considered fields since they do have a character separating them. We will need to specify to the sort/merge program what the field separator character is. We will also need to specify from which field the input or output keys are to come. This is done by placing the field number in parenthesis before the starting and ending key columns. For example, to specify the phone number in "names2" we would type "(3)1-7". This says to use columns one through seven of field three. The column numbers no longer refer to the entire input record but rather to the particular field specified. If no field separator character is specified (defaults to a null as in the examples above), sort/merge assumes that the input record is all one field. Note that the field number specification in a key defaults to one if there is not a parenthesis enclosed number. Thus in all the example sorts above, we were defaulting to field one since there was only one field.

Let's try sorting "names2". Suppose we wish to sort on the phone numbers and output the number followed by last name, comma, first name. We would begin by typing

```
++ sort names2
```

The first several prompts and responses should be exactly as before and are shown here as they would appear on your screen.

```
=== UniFLEX Sort Parameter Editor ===
```

```

Fixed or variable length records (f or v*)? v
EOR character or field count (default is EOR=$0D)?

```

At this point we receive a prompt for a field separator character. We wish to specify a comma and there are two ways to do that. You may specify the character as a hex value by prefacing the value with a dollar-sign ('\$') or as an ASCII character by prefacing the character with an apostrophe or single quote. For a comma the prompt and response would be

UniFLEX Sort/Merge User's Manual

Field separator character? \$2C

or alternatively,

Field separator character? ',

Enter one of these responses and continue as seen here

Output from key, input, or tag (k, i*, or t)? i

Enter input keys. Defaults to "a(1)1-10".

? (3)1-7

?

Enter output keys. Defaults to entire record.

? (3)1-3, '- ', (3)4-e, ' ', (1), \$2C, \$20, (2)

?

There are a couple of new concepts presented in these output keys. First is the key "(3)4-e". The 'e' in place of the ending column number represents the end of the field much as an asterisk represents the end of the input record. Thus this output key says to start outputting from column 4 of field 3 and continue outputting until a field separator is found (without outputting the field separator). The second new concept is the way in which fields one and two have been specified. You will note that the field number is present in parenthesis but that no starting or ending columns are specified. This is a special case of an output key which means to output the entire field. Thus an output key of "(2)" would be equivalent to "(2)1-e". This shorthand notation often saves much typing.

The rest of the prompts may be defaulted as before. When the sort is initiated, you should receive the following output.

```
422-1267 SOUTH, MARGARET
422-6002 SMITH, JERRY
422-6008 PERRY, ARNOLD
422-7264 SMITH, HAROLD
422-9105 FREEMAN, LINDA
457-2091 ABBOTT, JIM
457-5573 TYLER, HENRY
457-8339 WILSON, JIM
853-8227 LYON, WILLIAM
925-8411 CASWELL, AMY
```

To make another point, let's try sorting "names2" with respect to the last name. All prompts would be answered as before (set the field separator to a comma) until the input keys are requested. Now we know

we want field number one since that is where the last name is, and we know we want to start the key in column one of that field, but what should be the ending column. It might seem natural to specify "(1)1-e" to use column one through the end of the field but that is only allowed on output keys - not input keys. All input keys must be of the same length, so it is required that you specify that length by giving an ending column number. You should in fact specify an ending column that will fit the largest number one field found in any of the input records. If there are less actual columns in the field than you specify, the sort/merge program will "pad" the key with spaces to fill out the number of columns specified. In our case, we know that field number one is never longer than 10 characters, so we might specify "(1)1-10" or using the default field, simply "1-10". Note that you can always specify a key that is longer than necessary if you are unsure of the largest field. The only problem with this is that the larger the input keys, the more memory is required and the longer the sort will take. Enter the input key of "(1)1-10" and an output key of your choice (or default to outputting the entire record). The sort key that is built up of the input key for the first record would have 10 characters in it, the letters "WILSON" followed by four spaces. Executing the sort with the entire input records being output will yield the following.

```

ABBOTT,JIM,4572091
CASWELL,AMY,9258411
FREFMAN,LINDA,4229105
LYON,WILLIAM,8538227
PERRY,ARNOLD,4226008
SMITH,HAROLD,4227264
SMITH,JERRY,4226002
SOUTH,MARGARET,4221267
TYLER,HENRY,4575573
WILSON,JIM,4578339

```

We will find out later that sort/merge has an option which will print "run-time" messages as the sort is taking place. These are not errors, but informative messages as to what is occurring at different points in the sort since it can require considerable time to complete. If that option were selected (you will find how to select it later) in the above sample, you would see the following message at the end of the sort process:

```
Key padding was required
```

This is not an error message, just a report of the fact that padding was required.

The same type of padding will be done on output keys if necessary. There will, however, be no report of such as with the input key padding. For instance, an output key in the preceding example of "(2)1-15" would print the first name and then spaces until 15 columns had been printed.

3.5 SORTING WITH RIGHT OR LEFT JUSTIFICATION

Let us now reiterate somewhat on right and left justification. It should be obvious what right and left justification can be used for on output keys. It allows you to align columns of words or digits for a nice looking printout or to do simple formatting of an output file. However, the use of justification of input keys is not quite so obvious. Left justify on input keys will probably be rarely used. It would allow you to alphabetize some portion of a record which was not already justified, but this is probably a rare case. There is a good use for right justification of input keys. It relates to sorting numbers. Suppose we have a short file of dollar values that looks like this:

```
8.75
1.25
62.00
225.65
1.00
```

If we sorted this file with an input key of "1-10", the output would be ordered as follows.

```
1.00
1.25
225.65
62.65
8.75
```

This is due to the fact that sort/merge compares keys one column at a time from the left. If you examine the first column only, you will find that they are in order. If, however, we sorted the file with the input key right justified such as "r1-10", the result would be the proper numerical order. The keys which sort/merge would have built up would look like:

```
8.75
1.25
62.00
225.65
1.00
```

You can see that if these keys are sorted column at a time from the left, the correct order will result since the spaces are considered lower in value than any of the digits.

4.0 GENERAL USER SPECIFICATIONS AND SORT/MERGE DETAILS

There are several things which were touched on in the preceding tutorial which may be specified to the sort/merge module such as input and output keys. These specifications are the same for all five parameter supplying commands and are therefore elaborated on in this section. This section should be read prior to the sections which follow on the individual commands.

4.1 SPECIFYING INPUT RECORDS

As seen before, each input file is made up of "records" of information. These records are to be sorted into some logical order. Since these input records can vary in type and size, we must have some way of specifying to the sort/merge program where one record ends and another begins. There are two basic types of input records, "fixed length" and "variable length". Fixed length records are as the name implies records which will all be the same length. It is not necessary to have some character to mark the end of the records, but rather simply to specify how long the record is in number of characters. Variable length records do not have to all be of the same length. They are specified in one of two ways. The first is to specify some particular character which signals the end of the record. This character is called an "End Of Record" character. The second method is to specify some field count. In other words, the user would specify a count which would be the number of fields included in each record. These fields are signaled by an End of Field character as described before. Thus you may have a file which is made up of a large number of fields and split it into records by giving a field count.

Most files will probably be sorted as variable length files with an end of record (EOR) character. All five sort/merge commands default to this type with a carriage return (hex 0D) as the EOR. All our examples in the preceding section were done with this type of record. Note that the EOR character is never included in the input record that the sort/merge sees. It is in effect "swallowed up" as the records are read.

If we wanted, we could have sorted the "names" file as a fixed-length record file. The only problem would be that the carriage return would then NOT be swallowed up by sort/merge. The carriage return would have to be part of the record. If we wanted to do this, we would simply specify the decimal number of characters to be put into each record. In the case of "names" that number would be 28. We could have answered the prompts associated with record specification as follows.

Fixed or variable length records (f or v*)? f

Answering thusly with an 'f' yields the prompt for the record length which should be answered as shown

Record length? 28

The next prompt would be for the field separator character and would continue as before. The one problem with trying to sort the file as a fixed length file is that the carriage return is now part of the record. If you instruct the sort program to output the entire input record, either by defaulting or with an output key that includes all, the carriage return in the record will be output along with the carriage return which is always output at the end of an output record. This may not be the desired effect. There is a way around this problem, however. The carriage return character which is added to the end of every output record may be changed to any desired character or simply turned off. This is done in the special options section of "sort" as described later. Thus we could turn off the carriage return added to the end of our output records and only have the one that is part of the input record.

One point should be noted about input records. Any input record which is a null record (ie. a variable length record containing only the end of record character) is ALWAYS automatically deleted from the sort/merge process. There will be no indication of any such deletions to the user.

4.2 SPECIFYING INPUT KEYS

The sort tutorial section discusses the use and specification of input keys. This section is meant to elaborate somewhat on that discussion, stating the general format and limitations of input in more detail. A general syntax of a single input key specification would appear like:

<options>(<field #>)sss-eee

The "sss" represents the starting column number of the key within the field in use. If no field separator character has been defined, the entire input record is considered to be one field of number one. The "eee" represents the ending column of the key. These values must be between 1 and 250 inclusive. Note that "eee" may be smaller than "sss", in which case the key would be backward. In other words, the comparison of keys would take place one column at a time, beginning with column "eee" and continuing as necessary through column "sss". It is also permissible to have "sss" equal to "eee" in which case the input key would be only one character long. The starting and ending columns are ALWAYS required.

The <field #>, which must be enclosed in parenthesis as shown, is the number of the field from which the key should come. Fields are numbered starting with one and there may be up to 64 fields in each record. This portion of the key specification is optional and will default to field number one if omitted.

The <options> specification represents a list of options which can include the following:

- a ... Ascending sort order
- d ... Descending sort order
- l ... Left justify the key
- r ... Right justify the key

At most there should only be two, 'a' or 'd' and 'l' or 'r'. If conflicting options are specified (both ascending and descending or both left and right justify) the last one specified will take precedence. These option letters may be specified in any order, so long as they come before the field number and column specs. It is not required to give any options. If this is the case, the key defaults to ascending with neither right or left justify.

Up to 20 input keys may be specified and may come in any order. That is to say, the first key might come from the end of the record while the second key came from the start. Keys may also overlap each other's position in the input record.

4.3 SPECIFYING OUTPUT SOURCE

The sort does not actually sort the entire input records. It only sorts the input keys which were specified. These input keys have pointers to their parent record's location on the disk. When it comes time to output the input records in order, the sort program looks at which key is highest, finds out where it came from, and then re-reads that input record, writing it to the output as specified by the output keys. This means the "source" for the output is the input file. It is possible, however, to specify that the output data come from another location, namely the key itself. In the "sort" command, for example, we see:

Output from key, input, or tag (k, i*, or t)?

In all the previous examples, we selected the input as our output source. By simply typing a 'k', we can cause the output to come from the sort key itself. It is still possible to default to outputting the entire record (in this case the entire key) or to specify output keys with the columns now referring to the columns in the key rather than the input record. Note that there will rarely be fields to specify if outputting from the key. The advantage to outputting from the key is that it is considerably faster than outputting the input record since it is not necessary to re-read the input file as described above. There are two disadvantages to outputting from the key. First is that all the data from the input may not be in the key. Obviously when outputting from the key, only those portions of the input record which have been used as key data may be output. The second disadvantage is that if the upper case equal to lower case option is selected, all letters in the key will have been converted to upper case. This may or may not be acceptable depending on the situation.

The "tag" in the above prompt for output source refers to tag file output. A "tag" file is an output disk file which contains only the pointers to each key's parent input record. Each pointer is only five bytes, so we can maintain the necessary pointers for a sorted list in a very small space. This might be useful in the case where you wanted to keep several different ordered lists of a large file, but had limited disk space. These tag files do, however, require some external means of re-reading the input file in the order specified. More information on tag files is given in a later section.

4.4 SPECIFYING OUTPUT KEYS

Output key specification was covered in the preceding sections, but is repeated here in somewhat more detail. There are five types of output keys, each of which simply specifies some set of data to be sent to the output record.

- 1) The Full Key Spec... This is the most commonly used output key and is very similar to the input key discussed earlier. It has the general form:

`<justify>(<field #>)sss-eee`

The "sss" represents the starting column of the key within the field in use. It is a number from 1 to 250 inclusive. If no field separator character has been defined, the entire input record (or entire sort key depending on what was selected as output source) is considered to be field number one. The "eee" represents the ending column of the key. It can be a number between 1 and 250 inclusive or may be one of two special characters. The first is an asterisk ('*') which represents the ending column of the current input record (or of the sort key). Thus "1-*" as an output key would output the entire record. The second special character is the letter "e". It represents the end of the current field. Thus an output key of "1-e" would output all of field number one (the default field). Unlike input keys, "sss" may not be larger than "eee". They may, however, be equal. The <field #>, which must be enclosed in parenthesis as shown, is the number of the field from which the key should come. This portion of the output key specification is optional and will default to field number one if omitted. The <justify> portion of the output key specification is a one character option to either right justify the particular key (specified by the letter "r") or left justify (specified by the letter "l"). This portion is optional and defaults to neither left or right justified.

- 2) Entire Field Spec... This type of key is much like the previous type except that the field MUST be specified and no starting or ending columns are specified. The entire field is sent to the output record. For example, to output all of field 3 we could enter "(3)" as an output key which would be equivalent to "(3)1-e". This type of key simply saves typing over the first type.

- 3) Literal String... This type of output key allows the insertion of some constant string of characters (a literal) into the output record. The same string will be sent to each record output. The string of characters is specified by enclosing the characters in single quotes. Any printable ASCII characters may be included except for the single quote character itself.
- 4) Hexadecimal Value... Any single, 2 digit hexadecimal value may be inserted in the output record by specifying it with a preceding dollar-sign. For example, to insert a carriage return in the output simply type "\$OD". This allows any ASCII character to be output, printable or non-printable. If more than two hex digits are typed, only the last two are used with the others being ignored.
- 5) Horizontal Tab... Another type of output key is the horizontal tab which allows you to tab over to some particular column number in the output record with spaces used for padding. This key is specified as an at-sign followed by the decimal column number. For example, to tab over to column 25, enter "@25". If you are already past the column specified in a tab key, the tab key specification will be ignored.

The output keys may call data from the input record or sort key in any order. That is to say, data from field 4 may precede data from field 2. The only limit on the number of output keys is the amount of space reserved for the input and output keys. This should always be sufficient unless several very large literal string keys are specified.

4.5 SPECIFYING OUTPUT REDIRECTION

The UniFLEX Sort/Merge Package sends its sorted output to the standard output channel with which it was called. The standard output device is normally the Cterminal or CRT. Output may be sent to another device or disk file by using I/O redirection in the sort/merge command calling line. Note that even if standard I/O is redirected to some other device or file, all prompts, run-time messages, and error messages will be sent to the user's terminal via the standard error channel and will not appear in the actual sorted output. The user should consult documentation for the UniFLEX Operating System for further details of I/O redirection, but an example is given here:

```
++ sort names >sortednames
```

This would sort the file "names" and send the sorted output to a file called "sortednames" rather than the user's terminal.

4.6 SPECIFYING INPUT REDIRECTION

In most cases, the input to the sort/merge package (the input data to be sorted) will be in the form of a disk file. The name of this file or files is specified as part of the sort command line arguments. It is possible, however, to not specify an input file name in which case sort/merge will attempt to read the standard input channel. This implies that the input to sort/merge can be from a UniFLEX pipe. Piping input data into sort/merge does not make sense for two of the sort/merge commands, however. They are "sort" and "merge" because these two programs attempt to read the responses to their prompts from the standard input before the sort/merge operation ever takes place and would therefore be reading data as responses. Piping can be done to "csort", "psort", and "pmerge" and is a very useful technique. It allows the sort/merge package to act as a filter on data. An example will follow in the description of the "csort" command.

One important point should be made in regards to obtaining input data from redirected input rather than a named file. Since sort/merge must randomly re-read the input records in order when doing its output, it must copy the redirected input to a temporary scratch disk file before beginning the sort operation. The user should be aware of this fact and the increased time and disk space that go along with it.

4.7 LOCATION OF THE SORT/MERGE MODULE

The program that performs the actual sorting and merging is named "psort" and is one of the five programs that comprise the UniFLEX Sort/Merge Package. Since the other four sort or merge commands load and execute the "psort" command, it is imperative that "psort" be included in the active UniFLEX file system and that it be in a known directory so these commands can locate it. It is permissible to place "psort" into one of two directories. The first is the "/bin" directory where most used commands are generally kept. The second is the "/usr/bin" directory where less used commands are kept. When "sort", "merge", "csort", and "pmerge" attempt to load "psort", they first look for "/bin/psort" and if not found for "/usr/bin/psort". If the program is not found in either of these locations, the calling program will issue an error message and terminate.

All this implies that the "psort" file should be immediately copied from the UniFLEX Sort/Merge distribution disk into the "/bin" or "/usr/bin" directory of your active file system and left there permanently. It also implies that "psort" can NEVER be renamed.

4.8 TEMPORARY FILES

As explained in the section on Sorting Concepts, the Sort/Merge package must at times create temporary working files which are automatically deleted at the end of the sort/merge process. The user generally need not be concerned with these files but it is important to know about them should the system crash before these files have been deleted or in the case where another user or task has access to the directory in which they are created. Later in the manual you will find it is possible to instruct sort/merge as to the directory in which these temporary files should be created. By default they are created in the current working directory.

There are three kinds of temporary files created by sort/merge. The first is the parameter file. This is a file which contains all the necessary parameters for the sort/merge module to know how to perform a sort. In some cases this file is a permanent file that is not deleted upon completion of the sort or merge, but often it is a temporary file needed only for the duration of a single sort/merge operation. If it is a temporary file, it will be created with a name of the form "srtscrPF-TTTTT", where "TTTTT" represents the task id number of the task that created the file. The second type of temporary file is one which is a copy of the data input to sort/merge when input redirection is selected. This file is created with the name "srtscr00-TTTTT", where "TTTTT" is the task id number of the sort/merge task. The third type of temporary file is the type which is created when sort/merge does not have enough memory to carry out a sort in a single operation and must create sorted temporary files for later merging. These files are given names of the form "srtscrRR-TTTTT", where "TTTTT" represents the task id number as before and "RR" represents the number of the sort "run" which that file contains.

All these files should be deleted upon a normal termination of a sort or merge operation. If there is an abnormal termination for any reason and any of these files are left undeleted, they may be deleted by the user with the "kill" command in UniFLEX.

5.0 THE SORT COMMAND

The "sort" command is perhaps the easiest method of performing a sort operation as it prompts you for all the necessary parameters. This obviates the need to memorize what parameters must be supplied and how to supply them. The disadvantage is that for simple sorts which need few parameters you must still answer all prompts. This is simplified by accepting defaults for most prompts. Operation is as follows. Upon issuing a "sort" command, the "sort" module is loaded. This module interacts with the user, prompting for the necessary parameters. It is thus called a "parameter editor". When all parameters have been obtained, the user is allowed to save these parameters as a disk file if desired. Then he may exit the "sort" module (back to UniFLEX) or may continue with the sort. If elected to continue, the "sort" module will attempt to load the sort/merge module, "/bin/psort" or "/usr/bin/psort". If successful, the sort will then be performed.

5.1 GENERAL USE OF SORT

To initiate "sort", simply type a command of the general form:

```
++ sort file
    or
++ sort file1,file2,file3,...
```

Where "file" is a standard UniFLEX file specification for the file to be sorted. Note that more than one file may be sorted at a time. When one file has been completely read, it is closed and the next file is immediately opened for reading. The output file will consist of all the records of the specified input files. It is possible to simply type "sort" with no file specifications. This is useful only when the user wishes to edit a parameter file and not proceed with a sort operation.

The computer should respond to the "sort" command by typing

```
=== UniFLEX Sort Parameter Editor ===
```

This shows that the parameter editor has been entered and the prompts for sort parameters will follow. These prompts and the responses they require are described here. Note that the prompting may be restarted at any time by typing a 'CTRL Z'.

UnIFLEX Sort/Merge User's Manual

Fixed or variable length records (f or v*)?

Here the user may specify variable length input records by typing a 'v' or by defaulting with a carriage return. Alternatively he may specify fixed length records by typing an 'f'.

If fixed length records are chosen, the next prompt will be

Record length?

The record length in characters should be entered as a decimal number greater than zero.

If variable length records are chosen, the following prompt would be issued instead

EOR character or field count (Default is EOR=\$0D)?

The response to this prompt tells sort how to terminate a variable length input record. There are two possible responses. The first is to specify an End-Of-Record character as a hex value preceded by a dollar-sign or as a printable ASCII character preceded by a single quote. Notice that the default is a carriage return or 0D hex. The second possible response is a decimal field count. This tells sort to terminate an input record when the specified number of fields have been read.

Field separator character?

At this point a field separator character may be specified as a hex value preceded by a dollar-sign or as a printable ASCII character preceded by a single quote. If no field separator is desired, simply type a carriage return as the default is a null field separator character. Note that if a field count was specified as the input record terminator a field separator character will be required.

Output from key, input, or tag (k, i*, or t)?

This allows the user to specify the source for output data. The default is the input record. Typing a 'k' will cause the output data to come from the sort key itself. Typing a 't' will produce a tag file as described in a later section.

Enter input keys. Defaults to "a(1)1-10".

?

The input keys may be entered all on one line or on several lines. To put more than one key on a line, separate them with a comma or a space. When all desired keys have been entered on a line, hit a carriage return. The computer will respond with another question mark which is prompting for more keys. If you do not wish to enter more keys, the key prompt mode may be terminated by hitting a carriage return. See section 4.2 for details of input key specifications.

Enter output keys. Defaults to entire record.

?

Output keys are entered exactly like input keys with the capability to put all keys on one line or on multiple lines. See section 4.4 for details of output key specification.

Further options required (y or n*)?

At this point, the basic parameters necessary for a sort operation have been specified. There are, however, several other parameters or options which may be specified. If you need to set further options, type a 'y'. You will then be prompted for them as described in section 5.2. If no further options are required, type an 'n' or simply default with a carriage return.

At this point we see the message

=== Parameters are now set ===

This informs us that all the necessary parameters are set in the computer. We now have the opportunity to save these parameters as a file and then to exit or continue with the sort operation. The prompts for these functions follow.

Save parameter file on disk (y or n*)?

Typing a 'y' will allow the parameters just specified to be saved as a disk file. A prompt will be issued for the filename to be used. It is extremely important to note that the parameter file will be created by the name specified regardless of what is currently on the disk. If there is an existing file in the same directory by the same name given, it will be deleted and the parameter file written with no prompting or notification. If no permanent parameter file is needed, type an 'n' or a carriage return. In this case the parameters are saved in a temporary scratch file which will be automatically deleted when the sort operation is complete.

Exit or proceed with sort (e or s)?

At this point, the user has the option of exiting back to the operating system or continuing with the sort. Type an 'e' to exit or an 's' to continue with the sort. Note that there is no default. Hitting a carriage return will result in the prompt being re-issued.

If continuing with the sort, the sort/merge module will now be loaded and executed. Upon completion of the sort, control is transferred back to UniFLEX.

5.2 ADDITIONAL PARAMETER OPTIONS

There are several options in the sort which are less frequently needed. Instead of always prompting for them, the sort parameter editor keeps them as a separate group of prompts which will only be issued if desired. They are accessed by typing a 'y' in response to the prompt "Further options required (y or n*)?" as described above. When the last of these additional prompts has been answered, the message "=== Parameters are now set ===" will be issued and control will continue as described before.

There is one of these additional options that requires further elaboration. It is the Select/Exclude option and is described in section 5.3. The other prompts and required responses are as follows.

Directory name for temporary files?

This prompt allows the user to specify the directory into which any necessary temporary scratch files will be placed. The directory name (including path) may be up to 24 characters in length. The default (typing a carriage return) is the current directory.

Alternate collating sequence (y or n*)?

The sort/merge package normally does all sorting according to the ASCII collating sequence (see Appendix A). If the ASCII sequence is suitable, type an 'n' or simply a carriage return. If it is necessary to sort according to a different collating sequence, type a 'y'. You will then be prompted for the file name of the desired collating sequence file. For a description of the format of an alternate collating sequence file see section 12.0.

Treat lower case equivalent to upper (y or n*)?

In the ASCII coding scheme there is a different code for upper and lower case letters. This implies that an upper case 'E' would not sort equivalently to a lower case 'e'. In fact, the upper case characters are all lower in value than the lower case. Thus a 'Z' would be sorted before an 'a' if a normal ascending sort was performed. Typing a 'y' in response to this prompt will cause the sort program to treat upper case letters equivalent to lower case. In actuality, the sort keys are all converted to upper case. For this reason, if the output records come from the key instead of the input records, all letters would be upper case. Note that this feature is functional only if the ASCII character set is being used. If upper case SHOULD be sorted different from lower case, simply type an 'n' or a carriage return.

Delete records with blank sort keys (y* or n)?

Depending on how the key is specified and what the data contains, it is possible to end up with sort keys which are all spaces or blanks. Generally these keys are of no value since they contain no information. They simply take up space in memory and thus slow down the sort. These keys may be deleted from the sort operation by typing a 'y' or a carriage return. In effect, the input record

is hereby deleted. It still remains a part of the original data file, but no information from it is included in the output. This fact is alluded to at the end of a sort operation when a message is printed stating the number of records deleted. In some cases, a blank key may be meaningful and should not be deleted. If this is the case, type an 'n' in response to the prompt and the blank keys will be included.

Select/exclude option (y or n*)?

The sort/merge package has the ability to select or exclude certain records from the sort depending on their contents. If this option is not required, type an 'n' or a carriage return. If it is required, type a 'y'. Several related prompts will then be issued. The specifications required for this option are quite involved and thus a separate section of this manual has been devoted to describing them. It follows shortly as section 5.3.

EOR character for output records (Default=\$0D)?

The output key specifications allow the user to build up output records of any desired format. A built in function allows some end-of-record character to be appended to the data specified by the output keys. By default, this character is a carriage return (hex 0D). If a different character is desired, it may be specified here by typing a single quote followed by the ASCII character or by typing a dollar-sign followed by a two digit hex value for the character. If no end-of-record character is desired, type an 'n' for null (do not precede it with a single quote).

Print run-time messages (y or n*)?

The sort/merge package has several messages which it can print on the terminal during a sort operation to let the operator know just what is going on. These messages are explained in section 10.0. If the standard output is routed to a printer or to a disk file, these messages are not routed but rather are still printed on the terminal. To enable these messages, type a 'y'.

At this point, the message "=== Parameters are now set ===" and control will continue from this point as described in section 5.1.

5.3 SELECT/EXCLUDE OPTION

It is possible to have the sort package select only certain input records for the sort or to exclude certain input records. This is done by specifying a "select/exclude key". If this key matches the correct portion of the input record then that record will be selected or excluded. It is also possible to require that the key be greater than or less than the correct portion of the input record in order to be selected or excluded. This may sound a little confusing but can be cleared up by a couple of examples.

Suppose we wish to sort our file "names" according to last name and then first name but only want those people whose phone numbers have "422" as the exchange number (first 3 digits) to be in the output. In other words, we want to "select" the input records when the exchange is equal to 422. The sort parameters should be set up exactly as before with an input key of "1-20" to cause the sort to be done on last and first names. When the prompt "Further options required (y or n*)?" is received, type a 'y' so that we will receive the select/exclude option prompts. The prompts received may all be answered as desired or simply answered with a carriage return until you see the prompt:

Select/exclude option (y or n*)?

you should respond with a 'Y' which will cause the following prompt to be issued:

Select/exclude key spec?

This is a prompt for a SINGLE input key specification which tells what portion of the input record to which you wish to compare the select/exclude key. We want to compare to the exchange number, so you should respond with a "21-23" which are the columns containing the exchange. Now we receive the prompt:

Select or exclude (s* or e)?

This allows us to either select matching records or exclude them. We want to select all records with "422" as the exchange so type an 's' or simply a carriage return. We now see the prompt:

On key '<', '=', or '>' (Default is '=')?

Here we can specify that we want the records to be selected (or excluded had we so chosen) if the key is less than the specified portion of the input record ('<'), equal to it ('='), or greater than ('>'). We want to select if "422" is equal to columns "21-23" of the input record so type an equals sign ('=') or default with a carriage return. The final select/exclude prompt will now be issued:

Key string?

This is the prompt for the actual data which you want to select or

exclude on. We want to select on the exchange equal to "422" so simply type "422" followed by a carriage return.

At this point the prompts will continue as described in section 5.2. When the sort operation is complete, a run-time message will be printed informing you of the number of input records which were excluded from the output. If you were selecting records (as in our example) this would effectively be the number of records which were not selected.

The output from our example would look something like this (depending on what output key specifications you chose):

=== UniFLEX Sort/Merge ===

Sort Run 01 - 5 Records

| | | |
|---------|----------|---------|
| FREEMAN | LINDA | 4229105 |
| PERRY | ARNOLD | 4226008 |
| SMITH | HAROLD | 4227264 |
| SMITH | JERRY | 4226002 |
| SOUTH | MARGARET | 4221267 |

5 Records sorted
5 Records excluded

There are limitless possibilities to what can be accomplished with the Select/Exclude option. We could have easily excluded all records with an exchange of 422. We could have selected all records which had an exchange higher than 399 (i.e. 400 through 999). By doing multiple sort operations, we can even be more selective. Say for example we wish to sort all names which have a phone number with an exchange between 200 and 600 inclusive. First perform a sort selecting all records with an exchange of greater than 199 or excluding those which are less than 200. Then sort the output of this first sort selecting all records less than 601 or excluding all records greater than 600. We could even go a step further and exclude from this list all records which have a last name of "SMITH".

6.0 THE PSORT COMMAND

The "psort" command allows a user to supply all the necessary sort parameters in a named disk file. This file may be created through the use of the parameter editor (the "sort" command described in section 5.0). The "psort" command is the program which does the actual sorting and merging. All the other commands prepare a parameter file and then execute the "psort" command with that file. Directly performing a "psort" operation is especially convenient where one type of sort must be repeated on several files or repeated several times on one often changed file.

The syntax of this command is:

```

++ psort pfilename,inputfile
      or
++ psort pfilename,infile1,infile2,infile3,...
      or simply
++ psort pfilename

```

where "pfilename" represents the name of the parameter file to be used and "infile" represents the name of the input file. The first form shown loads the parameter file specified and proceeds to sort the specified input file. The second form shows that multiple input files may be specified. The third form shows that it is possible to not specify an input filename, in which case sort/merge will attempt to read the standard input device.

Several checks are made to ensure that the specified parameter file is actually a properly formatted parameter file. First, it must be exactly 496 bytes in length which is all that is required to contain all the parameters. Second, it must begin with the following four hex bytes: \$13, \$0F, \$12, \$14. These characters are to ensure the file is truly a parameter file. And finally, if called as "psort" itself and not by another command, the parameter file must be for a sort operation as opposed to a merge-only operation. Any parameter file produced by the "sort" command is for a sort operation, while a parameter file produced by a "merge" command is for a merge-only operation. If any of these conditions are not met, an error message will be issued and the sort operation will be aborted.

7.0 THE CSORT COMMAND

The "csort" command allows the user to specify the necessary sort parameters on the command line ("csort" stands for Command line Sort). This is very convenient for users who are proficient with the sort operation and don't need the prompting of the "sort" command or for those users who have an aversion for profuse prompting. The user simply types a single command line containing any desired parameters as shown below and hits a carriage return. The "csort" module writes the required parameter file, loads "psort", and immediately begins execution of the sort.

The general form of a "csort" command is as follows:

```
++ csort infile,"+input specs +output specs"
      or
++ csort infile1,infile2,...,infileN,"+input specs +output specs"
      or simply
++ csort "input specs +output specs"
```

where "infile" represents the input file to be sorted, "input specs" represents the input parameter specification as described below, and "output specs" represents the output parameter specifications also described below. It is important to note that the double-quotes (") and plus-signs (+) seen in the syntax above are actually part of the command. The input and output specs are enclosed in double-quotes so that the UniFLEX shell will not misinterpret any of the parameter specifications given. The plus-sign before the input specs tells "csort" that the following arguments are parameters and not filenames. The plus-sign before the output specs signals that the input specs are complete and the following arguments are output specifications. It is also important to note that the order shown above is critical. ALL input file names MUST come before the input and output specs which are enclosed in double-quotes. Everything past the command "csort" is optional. If the filename(s) are left out as in the third form above, "csort" will read the standard input channel for data to be sorted. If the input and output specs are left out, "csort" will assume the defaults listed below. If there are to be output specs but no input specs, both plus signs must be included. Several samples later in this section will help clarify the syntax of this command line.

Once a plus sign has been hit in the command line, "csort" begins looking for input specifications. These include input sort keys, record specifiers, etc. They may come in any desired order on the command line and are separated by either a space or a comma. The input specs may be any of the following:

UniFLEX Sort/Merge User's Manual

w=<directory name for temporary files>

A directory name into which all temporary scratch files should be placed may be specified. The name including path may be up to 24 characters in length. If this option is omitted, "csort" will default to the the current directory.

l=<decimal>

This is the input record length specification. <decimal> is the decimal length of the fixed length input records. It may be any number from 1 to 64K.

e=<hex or ASCII>

This is the End of Record character specifier for input records. <hex or ASCII> represents a hex character value preceded by a dollar-sign or an ASCII character preceded by a single quote.

c=<decimal>

This is the field count specifier for use when an input record is to be specified by a fixed number of fields. <decimal> represents the decimal number of fields desired.

f=<hex or ASCII>

This is the field separator character specification. <hex or ASCII> represents a hex character value preceded by a dollar-sign or an ASCII character preceded by a single quote. If this specification is left out, the field separator character will be a null, effectively disabling fielding.

u

If this character is entered on the command line, sort/merge will treat lower case characters equivalent to upper case characters so long as the ASCII character set is in use.

k

This character instructs the sort/merge package to "Keep" all records which have blank sort keys. Recall that normally these type records are deleted from the sort process. This specification turns off that deletion process.

standard input keys

Any standard input key may be included in the input specifications. For details on a standard input key specification see section 4.2.

A couple of comments on these input specifications are in order. First, the user may only specify one of 'l', 'e', or 'c' on a single command line. These specify when to terminate an input record and there can be only one method in use. If none of these three are given, "csort" will default to a carriage return (0D hex) as an End of Record character. Also note that if a field count is specified (c=<decimal>), then a field separator character must also be specified. If not an error message will result.

The input specifications may appear in any order but remember that the order of the input keys is significant. If there are conflicting specifications (ie. a 'w=/tmp' followed later by a 'w=/usr/joe'), the last one on the command line is used and the first is ignored.

If "csort" sees another plus sign while looking for input specifications, it will immediately terminate its search for input specifications and begin looking for output specifications. These output specifications are entered much as the input ones. That is to say they may come in any order and should be separated by a space or comma. Possible output specifications are as follows:

`o=<i, k, or t>`

This is the output record source specification (see section 4.3). The output may come from the input records by entering 'o=i', it may come from the sort key by typing 'o=k', or a tag file can be produced by typing 'o=t'. If none of these are entered, "csort" defaults to 'o=i'.

`e=<hex, ASCII, or n>`

This is the End of Output Record specification. <hex, ASCII, or n> represents a hex character value preceded by a dollar-sign, an ASCII character preceded by a single quote, or typing 'e=n' instructs sort/merge that the end of output record character is to be null, that is there will be NO character output at the end of output records. If this specification is left out, "csort" defaults to a carriage return as the end of output record character.

`m`

This is the message level specifier. Entering this character as an output specification will enable all run time messages during the sort/merge operation.

`standard output keys`

Any standard output key may be included in the output specifications. For details on a standard output key specification see section 4.4. Note that if no standard output keys are given, "csort" outputs the entire input record to the output record.

As before, the output specifications may appear in any order but the order of any output keys is significant. If there are conflicting specifications, the last one given is used with the first ones ignored.

Note that in order to make "csort" convenient to use, some of the features of sort/merge were not implemented. In particular, an alternate collating sequence is not allowed and the select/exclude option is not allowed. If the user needs these features, he has no alternative but to use the "sort" and "psort" commands.

A couple of sample "csort" command lines should help clarify its use.

1) The very first example of section 3.1 could be duplicated with the following command:

```
++ csort names
```

This uses defaults for all parameters much as was done by hitting carriage returns in section 3.1. Input keys default to "a(1)1-10" and the output defaults to outputting the entire input record as is.

2) The last example of section 3.3 could be duplicated with the command:

```
++ csort names "+1-20+r11-20,$20,1-10,'    phone no. ',21-23,'-',24-25"
```

Here we have specified an input key of "1-20" and several output keys as discussed in section 3.3.

3) The first example of section 3.4 could be duplicated by:

```
++ csort names2 "+f=$2C,(3)1-7+(3)1-3,'-',(3)4-e,' ',(1),$2C,$20,(2)"
```

Here we set the field separator to hex 2C (a comma) before giving the input key spec.

4) Following are some valid "csort" command lines which demonstrate placement of the input and output specs and output redirection.

```
++ csort junk1,junk2,junk3,"+e=$02,f=$0D,10-15"  
++ csort data "+10-20,1-9" >outfile  
++ csort test.dat "+100-115 +o=k,m,1-10"  
++ csort phone,"+u,k,31-40,1-10,21-30,+m" >directory
```

One point that should be noticed from the above samples is that the plus signs which separate input and output specs may or may not be preceded by a delimiter (space or comma) as desired by the user.

5) The following example shows how the standard input device can be read for the input data to be sorted. This is accomplished by omitting any filename to be sorted. In this example we will pipe the output of a UniFLEX "dir" command into a "csort" operation and sort the lines of output from the dir command according to the file size. To do this we must specify a "+b" option to "dir" (gives file sizes in bytes) and then sort on columns 15 thru 24 which is where the file size is located in the output lines from "dir". We will also specify a second sort key which is the filename so that if two files are the same size they will be alphabetized. Here is what we end up with.

```
++ dir +b ^ csort "+15-24,1-14"
```

Note that piping the output of the "dir" is a method of "redirecting input" to the following command which is our csort. The key "1-14" specifies the columns in which the the filenames are found.

8.0 THE MERGE COMMAND

To this point, most descriptions have centered around the sort operation. We now turn our attention to the merge-only type operation. Note that merging may take place in a sort but it is transparent to the operator. A "merge-only" operation means that no actual sorting is to be done. Instead, two or more files that are assumed to be already individually sorted are merged into one, ordered file. This is done by looking at the top record of each input file and selecting the lowest record (or highest depending on the order of the merge) to be sent to the output. The next record in that input file is then brought to the top for the next compare. Note that if the input files are not in sorted order the output file will not be in order. The "merge" command is much like the "sort" command in that it is essentially a "parameter editor" to setup the parameters for the sort/merge module.

To initiate a merge-only operation enter a command of the form:

```
merge file1,file2,file3,...
```

Any number of input files may be specified so long as they fit on the command line. It is also possible to simply type "merge" with no input file specifications. This is useful only when the user wishes to edit a parameter file and not proceed with the merge operation.

The computer should respond to the "merge" command line with:

```
=== UniFLEX Merge Parameter Editor ===
```

This shows that the parameter editor has been entered and the prompts for sort parameters will follow. The prompts which are issued are identical to those of the "sort" command of section 5.0 and require the same responses with two exceptions. The first is that the Select/Exclude option is not allowed in a merge-only operation. Thus there is no prompt for such in "merge". The second difference is that the final prompt is now "Exit or proceed with merge (e or m)?". This prompt may not be defaulted but must be answered with an 'e' or 'm'.

The remainder of the prompts are identical to those in "sort" and the reader is directed to section 5.1 for further descriptions. As in the "sort" command, a 'CTRL Z' will restart the prompts from the top.

9.0 THE PMERGE COMMAND

The "pmerge" command allows a user to supply all the necessary merge-only parameters in a named disk file. This file may be created through the use of a parameter editor (the "sort" command described in section 5.0 or the "merge" command in section 8.0). The "pmerge" command simply loads and executes the "psort" module which then loads the parameter file and proceeds with the merge. This type of merge operation is especially convenient where one type of merge operation must be often repeated. It is also valuable for performing sorts on several files using a parameter file sort and then merging these files with "pmerge" using the same parameter file.

The proper syntax for this command is:

```
++ pmerge pfilename,infile1,infile2,infile3,...
```

where "pfilename" represents the filename of the parameter file to be used for merging and "infile" represents the filename for the input files to the merge.

Two checks are made to ensure that the specified parameter file is actually a properly formatted parameter file. First, it must be exactly 496 bytes in length which is all that is required to contain all the parameters. Second, it must begin with the four hex characters: \$13, \$0F, \$12, \$14. If either of these conditions are not met, an error message will be issued and the merge operation will be aborted.

It should be noted that "pmerge" can make use of a parameter file which was prepared as a merge-only parameter file OR one prepared as a sort parameter file.

10.0 RUN-TIME MESSAGES

There are several run-time messages which may be printed out during the operation of the sort/merge package. These are not to be confused with error messages as they report no error. They simply report on the status of the sort/merge since the operation can require quite a lengthy period of time. These messages are printed by the "psort" module which performs the actual sorting and merging. One message is simply a header at the outset of a sort/merge operation, one is printed for each sort run, one for each merge pass, and up to four messages are printed upon completion of the operation to summarize just what occurred. Recall that a sort run is one memory buffer full of data which is sorted individually and saved as a temporary file while other runs are sorted. These messages are printed on the operator console only. They are output through the standard error output channel and are therefore not redirected should the standard output be redirected. Thus the operator may monitor the sort operation at the console while output is being sent to a disk file or to a printer.

By default, these run-time messages are disabled. They may be enabled if desired as explained in the separate parameter supplying commands, "sort", "csort", and "merge". If not enabled, absolutely no run-time messages will be issued.

The messages are listed and explained here:

1) === UniFLEX Sort/Merge ===

This is simply a header message to let the operator know that the sort/merge operation has successfully begun. It is the first function that sort/merge performs.

2) Sort Run xx - yy Records

This message is printed for every sort run required. "xx" represents the sequential number of the run in progress while "yy" represents the number of records which are contained in that run. This line is actually printed in two parts. The sort run number is printed before any work has been initiated on the particular run. The number of records in the run is printed after the run has been read into the computer's memory. Thus you will see the sort run message, a pause while the records are being read from disk into memory, the number of records message, and then another pause while the run is sorted and written out to a temporary file.

3) Merge Pass xx - yy Runs

This message is printed for every merge pass required. This will generally be only one pass. "xx" represents the sequential number of the merge pass in progress while "yy" represents the number of runs being merged in that pass. This line is all printed at once before the merge pass has been initiated.

4) xx Records sorted

This message is printed at the completion of a sort operation. "xx" represents the total number of records that have been sorted and sent to the output. Note that this may not be equal to the number of records in the input file due to deleted records and excluded records.

5) xx Records merged

This message is printed at the completion of a merge-only operation. "xx" represents the total number of records that have been merged to form the output file. Note that this may not be equal to the total number of records in the input files due to deleted records.

6) xx Records deleted

This message is printed at the completion of a sort or merge-only operation. "xx" represents the number of records which were deleted from the process due to a blank sort key. This message is only printed if there were records deleted.

7) xx Records excluded

This message is printed at the completion of a sort operation if any records were excluded (or not selected) through use of the Select/Exclude option. If the Exclude option was set, "xx" represents the number of records which were excluded during the sort. If the Select option was set, "xx" represents the number of records which were NOT SELECTED or effectively excluded.

8) Key padding required

This message is printed at the completion of a sort or merge-only operation if input key padding was required. This does not denote an error condition but simply informs the user that the padding was required. If this should not have occurred, it is up to the user to remedy the situation.

11.0 ERROR MESSAGES

There are several error messages associated with the five programs which comprise the sort/merge package. Many of them are common between routines and will therefore be grouped together in this section to avoid redundancy. We will place the messages in three groups: those from "sort" and "merge", those from "csort", and those from "psort" and "pmerge". The first two groups are the messages which can come from those parameter supplying routines before the "psort" module, which does the actual sorting and merging, is loaded. After the sort or merge has begun, the error messages listed under the third group may be received as well.

11.1 ERROR MESSAGES FROM SORT AND MERGE

In general, the error messages associated with these two modules are of an interactive, non-fatal nature. In other words, they inform you of an error and allow you to re-enter the corrected data. They do not cause the program to abort. The messages are as follows:

- 1) Non-zero field count requires a field separator
This message is issued if the input records were specified by a field count and then no field separator character is specified. In order to count fields, there obviously must be some way of separating the input record into fields. The prompt will be re-issued.
- 2) Illegal key specification
An error was found in one of the keys in the line just entered. Any keys already entered are discarded and a new prompt is issued.
- 3) Too many keys. Re-enter.
There is a 256 byte buffer reserved for key information. Each input key requires 4 bytes and except for literal strings each output key requires 3 bytes. The number of input keys is limited to 20 while the number of output keys is limited by the amount of space left in the buffer. Literal strings as output keys are stored in the key buffer and thus long strings can rapidly eat up the available space. The 256 bytes should be sufficient, however, for most any application. If this message is issued, any existing keys will be discarded and a new prompt issued.
- 4) Can't create parameter file
UnifLEX detected an error in attempting to create the parameter file.
- 5) Error writing parameter file
UnifLEX encountered a write error while writing the parameter file.

- 6) Device filled writing parameter file
The disk device onto which the parameter file was being written ran out of free space.
- 7) Can't find 'psort' module
Once the parameters have been finalized, "sort" or "merge" attempts to load the actual sort/merge program called "psort". It first attempts to find it in the "/bin" directory ("/bin/psort") and if not there looks in the "/usr/bin" directory ("/usr/bin/psort"). If still not found, this message is issued and the program is aborted, returning control to UniFLEX.
- 8) Unable to execute 'psort' module
If "sort" or "merge" was able to find the "psort" module, but unable to execute it, this message is issued and control returned to UniFLEX. Reasons for not being able to execute the file include a missing directory and lack of permissions on the file.

11.2 ERROR MESSAGES FROM CSORT

Errors in "csort" are generally of a fatal nature in that they cause execution of the sort to be terminated and return control to UniFLEX.

- 1) Unrecognizable input spec
"csort" has found an unrecognizable input specification. In other words, "csort" was looking for an input specification such as an input key, the temporary directory name, or a record terminator spec, but instead found some unrecognizable item. Note that a recognizable input spec which is in error will not produce this message but rather a more specific message as seen later.
- 2) Unrecognizable output spec
"csort" has found an unrecognizable output specification. In other words, "csort" was looking for an output specification such as an output key, the output source, or an end of output record character spec, but instead found some unrecognizable item. Note that a recognizable output spec which is in error will not produce this message but rather a more specific message as seen later.
- 3) Input key error
An error was found in one of the input keys specified.
- 4) Temporary directory name too long
A directory was specified for temporary files that was too many characters in length. "csort" will only accept a name of up to 24 characters.
- 5) Illegal record length
An invalid number was given as the record length. The value must be decimal and must not be zero.

- 6) Illegal end of record character
An invalid end of record character was specified. The EOR character must be specified as a hex value preceded by a dollar-sign or an ASCII character preceded by a single quote.
- 7) Illegal field count
The field count specified is invalid. It must be a decimal number between 1 and 255 inclusive.
- 8) Only one of l, e, or c may be specified
This message is issued if more than one of the three record defining specs was found. It is only possible to have one of the three in a single command.
- 9) Illegal field separator
The field separator character was specified incorrectly. It must be either a hex value preceded by a dollar-sign or an ASCII character preceded by a single quote.
- 10) Non-zero field count requires a field separator
This message is issued if a field count has been specified but no field separator. In order to have a field count, there must be a field separator character.
- 11) Output key error
An error was found in one of the output keys specified.
- 12) Illegal output source specified
The output source may be specified only as 'i', 'k', or 't' for input, key, or tag.
- 13) Illegal end of output record character
An invalid end of output record character was specified. It must be a hex value preceded by a dollar-sign, an ASCII character preceded by a single quote, or the letter 'n' which signifies a null end of output record character (no EOR character appended to record).
- 14) Too many keys specified
Sort/Merge reserves 256 bytes for input and output key specifications. The user is limited to 20 input keys and as many output keys as will fit in the remainder of the 256 bytes not used by input keys. This message is issued if more than 20 input keys were specified or if the input and output keys overflowed this 256 byte limit. The only recourse is to lower the number of specified keys.
- 15) Equals sign required in spec
As seen in section 7.0, many of the input and output specs are a single letter, followed by an equals sign, followed by the actual parameter. If the equals sign is omitted from one of these specs, this error message will result.

16) Syntax error

This message is issued if there is a syntactical error in the command line such as more than two plus signs.

17) Can't create parameter file

UniFLEX detected an error in attempting to create the parameter file.

18) Error writing parameter file

UniFLEX encountered a write error while writing the parameter file.

19) Device filled writing parameter file

The disk device onto which the parameter file was being written ran out of free space.

20) Can't find 'psort' module

Once the parameters have been finalized, "csort" attempts to load the actual sort/merge program called "psort". It first attempts to find it in the "/bin" directory ("/bin/psort") and if not there looks in the "/usr/bin" directory ("/usr/bin/psort"). If still not found, this message is issued and the program is aborted, returning control to UniFLEX.

21) Unable to execute 'psort' module

If "csort" was able to find the "psort" module, but unable to execute it, this message is issued and control returned to UniFLEX. Reasons for not being able to execute the file include a missing directory and lack of permissions on the file.

11.3 ERROR MESSAGES FROM PSORT AND PMERGE

These errors are generally fatal in that they cause the execution of sort/merge to be terminated and control returned to UniFLEX. Note that since all sort/merge commands execute the "psort" module after they have prepared a parameter file, these messages can be received from any sort or merge operation.

1) Illegally called

"psort" must be called with a parameter file specified in the argument list. If there are no arguments (no parameter file) this message is delivered.

2) Illegal input file or device

This message is issued if the input to sort/merge is from a block device itself, a directory, or any other illegal device.

3) Illegal parameter file

There are two possible reasons for this message. First is that the specified parameter file is not of the correct length. The parameter file must be exactly 496 bytes in length to prevent invalid files from being loaded. Second, there is a special 4 byte header on parameter files which must be matched to ensure the specified file is truly a parameter file. Check to be sure you have specified the correct file.

4) Illegal sequence file

An alternate collating sequence file was specified that was not the correct length of exactly 256 bytes.

5) Too many files

This message can only be received during a merge-only operation. It is caused by an attempt to merge too many files. Note that in a merge-only operation, all merging must be done in one pass. The only solution to this problem is instead of a single large merge, merge the files in several smaller groups which can then be merged.

6) Merge requires at least 2 files

This message can only be received during a merge-only operation. It is caused if less than two input files are specified.

7) Sort key too long

The total length of all the input keys specified must not exceed 250 bytes. If it does, this error message is issued and execution is terminated.

8) Too many work files

Sort/Merge allows a maximum of 99 temporary work files. This should never be a limitation, but if the number is exceeded, this error message is issued.

9) Parameter file is merge-only

This message can only be issued if the calling command was "psort". It is issued if the parameter file specified to "psort" was prepared using "merge". This arrangement is not permissible.

10) Sort/Merge terminated by INTERRUPT!

Sort/Merge catches interrupts caused by hitting a 'CTRL C' or a 'CTRL \'. This is so that any temporary work files can be deleted before returning to UniFLEX.

11) Read error in 'filename'

UniFLEX has encountered an error in reading the specified file.

12) Write error in 'filename'

UniFLEX has encountered an error in writing the specified file.

- 13) Device full writing 'filename'
As the specified file was being written, the disk on which it resides became full (no more free space).
- 14) Bad file error with 'filename'
A "Bad File" error was encountered with the specified file. See UniFLEX documentation for a more detailed description of this error.
- 15) Can't find 'filename'
UniFLEX was unable to carry out the requested disk operation because it was unable to locate the specified file.
- 16) Missing directory in 'filename'
One of the directory path names in the file specification is non-existent or lacks proper permission.
- 17) No permission for 'filename'
An attempt was made to access the specified file without proper permission for that file.
- 18) File is a directory: 'filename'
The specified file is a directory but should be a true file for the requested operation.
- 19) Can't find 'psort' module
This message is only issued from "pmerge". "pmerge" attempts to load "psort" and execute it with the specified parameter file. It first attempts to find it in the "/bin" directory ("/bin/psort") and if not there looks in the "/usr/bin" directory ("/usr/bin/psort"). If still not found, this message is issued and the program is aborted, returning control to UniFLEX.
- 20) Unable to execute 'psort' module
If "pmerge" was able to find the "psort" module, but unable to execute it, this message is issued and control returned to UniFLEX. Reasons for not being able to execute the file include a missing directory and lack of permissions on the file.

12.0 ALTERNATE COLLATING SEQUENCE FILE

Sort/Merge generally performs all sorting and merging according to the ASCII sequence of characters (see appendix A). Thus for example, a '3' is higher in the sequence than a 'P' which is higher than a 'g'. It is possible to sort according to some other sequence or to use a different code than ASCII. This is done thru the use of an "alternate collating sequence file" which is simply a named disk file which contains all the characters listed in the desired order (each character is represented by a single byte). Thus we could use the same ASCII code (ie. '3' = 33 hex, 'P' = 50 hex, and 'g' = 67 hex) but place the characters in a different sequence in the alternate collating sequence file so that the collating order of these three characters might be 'P', 'g', and '3' or possibly 'g', 'P', and '3'. It is also possible to specify a totally different character set such as EBCDIC by simply specifying the proper character representations in the desired order.

A proper alternate collating sequence file MUST be EXACTLY 256 characters in length. Any other size file will cause an error in sort/merge.

There is no particular method provided for creating an alternate collating sequence file. It is up to the user to create his own. It is a relatively simple procedure in UniFLEX BASIC, however, and a sample program is listed here which when run will produce a valid alternate sequence file. This file uses the ASCII code, but sorts spaces first, followed by upper case, followed by lower case, followed by numbers, followed by the graphic characters and control characters.

```

10 rem Alternate collating sequence file generator
20 width 0
30 open new "altseq" as 1
40 print #1,chr$(32);
50 for i%=65 to 90:print #1,chr$(i%);:next i%
60 for i%=97 to 122:print #1,chr$(i%);:next i%
70 for i%=48 to 57:print #1,chr$(i%);:next i%
80 for i%=33 to 47:print #1,chr$(i%);:next i%
90 for i%=58 to 64:print #1,chr$(i%);:next i%
100 for i%=91 to 96:print #1,chr$(i%);:next i%
110 for i%=123 to 127:print #1,chr$(i%);:next i%
120 for i%=0 to 31:print #1,chr$(i%);:next i%
130 for i%=128 to 255:print #1,chr$(i%);:next i%
140 close 1
150 end

```


13.0 TAG FILE OUTPUT

There is a special type of output from the sort/merge package that warrants our attention. That is the tag file. The "tag file" output is a special output file which contains only the sorted pointers back to the original file. Thus each output record has no data from the input record, but has five pointer bytes to the file character position of the start of the corresponding input record. The first byte is a "file number". This is the ordinal number of the input file on the calling line. If there is only one input file, the first byte of all tags would be a "1". The remaining four bytes of each tag is the file character position for the particular file numbered in the first byte. Thus to find the input record for a particular tag, we would find which input file the record came from by examining the first byte of the tag and the list of input files and then use the remaining four bytes of the tag as a seek address from the beginning of that file.

The purpose of the tag file is to provide a very succinct form of storing information needed to produce a sorted file. This might be necessary where the file being sorted is very large and there is little disk space remaining or where it is desired to maintain several sorted versions of one database. Keeping several of these tag files would be much more efficient than the wasteful and redundant method of keeping several sorted copies of the database itself.

Tag files alone are quite useless. There must be some user written program to make use of the tag files and the original database as desired to print out a sorted list or perform some other function. The Sort/Merge package has no provisions for using the tag files produced. That is left up to the user for his own particular application.

14.0 PARAMETER FILE DESCRIPTION

All sorts and merges are actually performed as though they were a "psort". "sort", "merge", "csort", and "pmerge" all prepare a parameter file and then execute "psort" passing the parameter file to it. This parameter file contains all the necessary parameters for the actual sorting and merging to take place. It is fixed in length at 496 bytes and has a specific format as described below. Following the file layout, further descriptions of each parameter will be given. At the end is a description of how a user can prepare his own parameter file and call "psort".

14.1 PARAMETER FILE LAYOUT

| Location | Name | Description |
|-----------|--------|--------------------------------------|
| 0000-0003 | HEADER | Parameter file header |
| 0004-0005 | OSPEC | Output key specs pointer |
| 0006 | EOR | End Of Record character for input |
| 0007 | EOF | End Of Field character |
| 0008 | EOOR | End Of Output Record character |
| 0009 | GROUP | Field count or group count |
| 000A-000B | RLNGTH | Fixed length record length |
| 000C | LOWUP | Treat lower case equal to upper flag |
| 000D | FROMKY | Output from key designator |
| 000E | IDXTAG | Tag file output flag |
| 000F | OUTALL | Output entire input record flag |
| 0010 | MSGVL | Run-time message level |
| 0011 | DELNUL | Delete records w/ null keys flag |
| 0012 | SELXCL | Select/Exclude information |
| 0013-0017 | SESPEC | Select/Exclude key spec |
| 0018 | MGONLY | Merge-only flag |
| 0019 | DELPRM | Delete parameter file flag |
| 001A-0029 | | Reserved for future use |
| 002A-0042 | WRKDIR | Directory name for temporary files |
| 0043-006F | ALTSEQ | Alternate sequence file name |
| 0070-016F | ISPEC | Input and output specs |
| 0170-01EF | SEDATA | Select/Exclude string data |

UniFLEX Sort/Merge User's Manual

The following descriptions give the necessary details on what each parameter represents and what values it may take on.

- HEADER This four byte field is a fixed header to ensure that a valid parameter file has been loaded. The four bytes should be hex \$13, \$0F, \$12, and \$14 respectively.
- OSPEC These two bytes hold an offset value from the beginning of the input key specifications to the first byte of the output key specifications. The output key specifications should be placed directly after the input specs.
- EOR This is the End of Record character for input records. If input records are to be fixed length or if a field count is specified, this byte should be set to zero. Note that hex 00 may not be used as an actual EOR character.
- EOF This is the End of Field character used to separate input records into fields. If no fields are desired, this byte should be cleared to zero. This implies that 00 cannot be a true field character. This byte must be non-zero if a field or group count is set.
- EORR This is the End of Output Record character which is added onto the end of every output record. One exception is when this byte is zero, no EORR character is added to the output records.
- GROUP This byte holds the group count or field count. It is the number of fields required to make up one input record. If input records are specified by an EOR or a fixed record length, this byte should be cleared to zero.
- RLNGTH This is a two byte value containing the length of fixed length input records in binary. If input records are specified by an EOR or by a field count, these two bytes should be zero.
- LOWUP This flag may be set non-zero to cause lower case characters to be sorted equivalent to upper case characters. Note that this is only valid if using the ASCII character set.
- FROMKY This flag should be set non-zero to cause the data for the output records to come from the sort keys themselves.
- IDXTAG This byte signals the output file to be a tag file. The high order bit (bit 7) should be set. If not a tag file, this byte should be zero.

- OUTALL This byte may be set non-zero to indicate that the entire input record should be sent to the output record. If this is the case, sort/merge will ignore any output key specs and the output file will be a re-arranged copy of the input file.

- MSGVL This byte may be set non-zero to enable the printing of all run-time messages.

- DELNUL This byte may be set non-zero to cause any records with null sort keys to be automatically deleted from the sort operation.

- SLEXCL A non-zero value in this byte signals a select/exclude operation. The high order bit (bit 7) denotes select if cleared or exclude if set. The low order bits are set to 1, 2, or 3 to denote equal, greater than, or less than respectively.

- SESPEC These four bytes hold the select/exclude key specification. It is a single, standard input key spec as described below under the ISPEC description.

- MGONLY This flag may be set non-zero to signal a merge-only operation. If set, the entire sort operation will be skipped and the input files named will be merged.

- DELPRM If this byte is set, sort/merge will delete the parameter file as soon as it has read it into memory. To prevent automatic deletion of the parameter file, clear this byte to zero.

- WRKDIR These 25 bytes contain the specification of a directory into which sort/merge will place all temporary work files. The name string must be terminated with a zero byte. To default to the current directory, be sure the first byte of this area is a zero.

- ALTSEQ These 45 bytes contain the specification of an alternate collating sequence file. The file name should be terminated with a zero byte. If there is to be no alternate collating sequence, the first byte of this area must be equal to zero.

- ISPEC This is a 256 byte area in which all input and output key specifications are stored. The input keys are stored first followed by the output keys. The OSPEC value above points to the address in this area where the output keys start. Since the format of the input and output key specs are different, we will describe them separately.

INPUT KEY SPEC FORMAT

Each input key requires exactly four bytes of storage. The keys are stored adjacent in memory starting at 0100 hex. The input keys MUST be terminated by a zero byte immediately following the last input key. This appears to sort/merge like a key with zero length thereby terminating the fetching of input keys. The four bytes contain the following information.

BYTE 1: Length of key in bytes (1-250)
 BYTE 2: High order bit (bit 7) = 0 if forward
 = 1 if backward
 Low order bit (bit 0) = 0 if ascending
 = 1 if descending
 BYTE 3: High order bit (bit 7) = 1 if left justified
 Next high bit (bit 6) = 1 if right justified
 Lower bits (bits 5-0) = (field number)-1
 BYTE 4: Starting column number of key

OUTPUT KEY SPEC FORMAT

Output keys can be of variable lengths. The most common is three bytes in length and is essentially like bytes 1, 3, and 4 of an input key spec. The output keys immediately follow the zero byte after the input keys. OSPEC should be set to point to the first byte of the output key specs. As with the input specs, the output key specs are terminated by a following zero byte which looks to sort/merge like a key length of zero. A normal, three byte output key spec is setup as follows.

BYTE 1: Number of columns in key (1-250 decimal)
 BYTE 2: High order bit (bit 7) = 1 if left justified
 Next high bit (bit 6) = 1 if right justified
 Lower bits (bits 5-0) = (field number)-1
 BYTE 3: Starting column number of key

The next type of key is very much like the one just shown but specifies outputting to the end of the field or to the end of the record. This type key is always 3 bytes and is signalled by the first byte being FF hex for outputting to end of record or FE hex for outputting to end of field. Bytes 2 and 3 are just as before.

The next type of output key to describe is a string literal or hexadecimal value. This key must contain the string or a byte containing the hex value. Since the string can be any length, this type of key must also be of varying length. It is in fact of length N+2 where 'N' represents the number of characters in the string or is equal to one if a hexadecimal value. The format of this type key is as follows.

BYTE 1: Always equal to FD hex.
 BYTE 2: Number of characters in string (1 if a hex value)
 BYTES 3 thru (N+2): Actual string or hexadecimal value

The final type of output key spec is a tab. It is always two bytes in length and is as follows.

BYTE 1: Always equal to FC hex.
 BYTE 2: Column number to which to tab (1 to 255).

SEDATA This space holds the actual string data for select/exclude comparison if that option is in use. The string may be up to 128 bytes in length. If the string is not 128 bytes in length, the remainder of this area should be filled with space characters or 20 hex.

14.2 CALLING SORT/MERGE FROM OTHER PROGRAMS

In normal use, the parameter file will be created by using the "sort" or "merge" command or a temporary parameter file may be created by those commands or "csort". It is possible, however, for the user to prepare his own parameter file according to the above description. This could be done by a BASIC program, a Pascal program, an assembler program, or just about any means you desire. For example, you may wish to perform a sort from a BASIC program and have that program prompt for the sort parameters rather than require the user to directly use "sort" or "csort". Once the BASIC had prepared a parameter file according to the above rules, it might call up the sort/merge module ("psort") with a command like:

```
340 exec,"psort pfile datafile >sortedfile"
```

where "pfile" is the parameter file that BASIC created, "datafile" is the input file to be sorted, and "sortedfile" is the file which will be created to hold the sorted output.

Once the parameter file has been created, it is rather simple to call sort/merge from an assembler language program also. It simply requires a UniFLEX "exec" of "psort". The arguments for the exec should be setup so that the first argument is the calling program name, the second argument is the parameter file name, and the third argument is the name of the input file to be sorted. If there is more than one file to be sorted, the additional names become arguments four, five, etc. It is also important to note that sort/merge makes use of the standard input and output files or devices. Whatever the standard output file is defined as upon performing the exec to sort/merge is where the sorted output will be sent. The standard error output file is where all run-time messages and error messages are sent. The standard input file is read as the input data for sorting if there are no input file arguments (ie. only the command argument and the parameter filename argument).

15.0 APPENDIX A - ASCII CHARACTER SET

| HEX | CHARACTER | NAME |
|-----|-----------|--------------------------|
| 00 | CTRL @ | NUL |
| 01 | CTRL A | SOH |
| 02 | CTRL B | STX |
| 03 | CTRL C | ETX |
| 04 | CTRL D | EOT |
| 05 | CTRL E | ENQ |
| 06 | CTRL F | ACK |
| 07 | CTRL G | BEL,Bell |
| 08 | CTRL H | BS,Backspace |
| 09 | CTRL I | HT,Horizontal Tab |
| 0A | CTRL J | LF,Line Feed |
| 0B | CTRL K | VT,Vertical Tab |
| 0C | CTRL L | FF,Form Feed |
| 0D | CTRL M | CR, Carriage Return |
| 0E | CTRL N | SO |
| 0F | CTRL O | SI |
| 10 | CTRL P | DLE |
| 11 | CTRL Q | DC1 |
| 12 | CTRL R | DC2 |
| 13 | CTRL S | DC3 |
| 14 | CTRL T | DC4 |
| 15 | CTRL U | NAK |
| 16 | CTRL V | SYN |
| 17 | CTRL W | ETB |
| 18 | CTRL X | CAN, Cancel |
| 19 | CTRL Y | EM |
| 1A | CTRL Z | SUB |
| 1B | CTRL [| ESC, Escape |
| 1C | CTRL \ | FS |
| 1D | CTRL] | GS |
| 1E | CTRL ^ | RS |
| 1F | CTRL _ | US |
| 20 | | Space, Blank |
| 21 | ! | |
| 22 | " | |
| 23 | # | |
| 24 | \$ | |
| 25 | % | |
| 26 | & | |
| 27 | ' | Apostrophe, Single Quote |
| 28 | (| |
| 29 |) | |
| 2A | * | |
| 2B | + | |
| 2C | , | Comma |
| 2D | - | Minus |
| 2E | . | Period |
| 2F | / | |

UniFLEX Sort/Merge User's Manual

| <u>HEX</u> | <u>CHARACTER</u> | <u>NAME</u> |
|------------|------------------|--------------|
| 30 | 0 | Number zero |
| 31 | 1 | Number one |
| 32 | 2 | |
| 33 | 3 | |
| 34 | 4 | |
| 35 | 5 | |
| 36 | 6 | |
| 37 | 7 | |
| 38 | 8 | |
| 39 | 9 | |
| 3A | : | |
| 3B | ; | |
| 3C | < | Less Than |
| 3D | = | |
| 3E | > | Greater Than |
| 3F | ? | |
| 40 | @ | At Sign |
| 41 | A | |
| 42 | B | |
| 43 | C | |
| 44 | D | |
| 45 | E | |
| 46 | F | |
| 47 | G | |
| 48 | H | |
| 49 | I | Letter I |
| 4A | J | |
| 4B | K | |
| 4C | L | |
| 4D | M | |
| 4E | N | |
| 4F | O | Letter O |
| 50 | P | |
| 51 | Q | |
| 52 | R | |
| 53 | S | |
| 54 | T | |
| 55 | U | |
| 56 | V | |
| 57 | W | |
| 58 | X | |
| 59 | Y | |
| 5A | Z | |
| 5B | [| |
| 5C | \ | |
| 5D |] | |
| 5E | ^ | Up Arrow |
| 5F | _ | Underscore |

| <u>HEX</u> | <u>CHARACTER</u> | <u>NAME</u> |
|------------|------------------|---------------------|
| 60 | ` | Accent Grave |
| 61 | a | |
| 62 | b | |
| 63 | c | |
| 64 | d | |
| 65 | e | |
| 66 | f | |
| 67 | g | |
| 68 | h | |
| 69 | i | |
| 6A | j | |
| 6B | k | |
| 6C | l | |
| 6D | m | |
| 6E | n | |
| 6F | o | |
| 70 | p | |
| 71 | q | |
| 72 | r | |
| 73 | s | |
| 74 | t | |
| 75 | u | |
| 76 | v | |
| 77 | w | |
| 78 | x | |
| 79 | y | |
| 7A | z | |
| 7B | { | |
| 7C | | Vertical Slash |
| 7D | } | Alt Mode |
| 7E | ~ | (Alt Mode) |
| 7F | | DEL, Delete, Rubout |

NOTES:
